


Sum-of-Squares Lower Bounds for **Sparse** Independent Set



Chris Jones (University of Chicago) ← on the postdoc market

Aaron Potechin (University of Chicago)

Goutham Rajendran (University of Chicago)

Madhur Tulsiani (Toyota Technological Institute at Chicago)

Jeff Xu (Carnegie Mellon University)

Outline



The Sum-of-Squares algorithm and Independent Set



Pseudocalibration doesn't work



Matrix-valued Fourier analysis



Conclusion

The Sum-of-Squares algorithm and Independent Set

The Sum-of-Squares (SoS) Algorithm

SoS is a super strong **semidefinite programming (SDP)** algorithm

- SDP = linear program over the set of positive semidefinite (PSD) matrices

SoS can be used to optimize a polynomial subject to polynomial constraints

SoS consists of complete local reasoning (=Sherali-Adams) and PSD-ness

SoS captures our best algorithms for Max Cut [GW'95], Sparsest Cut [ARV'04], Unique Games [BBKSS'21], Tensor PCA [HSS'15], robust linear regression [BP'21], ...

- PSD-ness is critically important for all of these applications
- PSD-ness is also what makes proving lower bounds against SoS so hard

[Goemans-Williamson '95] [Arora-Rao-Vazirani '04]
[Bafna-Barak-Kothari-Schramm-Steurer '21] [Hopkins-Shi-Steurer '15]
[Bakshi-Prasad '21]

The Sum-of-Squares (SoS) Algorithm

SoS_D is actually a hierarchy of convex relaxations, one for each **degree** D . As D gets larger, the convex program gets larger and larger (slower runtime), but the value of SoS_D approaches the true value.

$$\text{val}(\text{SoS}_2) \geq \text{val}(\text{SoS}_4) \geq \text{val}(\text{SoS}_6) \geq \dots \geq \text{val}(\text{SoS}_{2n-2}) \geq \text{val}(\text{SoS}_{2n}) = \alpha(G)$$

Lemma. SoS_D can be expressed as a semidefinite program of size $n^{O(D)}$

$D = O(1)$ corresponds to polynomial time

$D = 2n$ solves the problem exactly in exponential time (for n Boolean variables)

Main Question: how large must D be before $\text{val}(\text{SoS}_D) \approx \alpha(G)$?

How to prove an SoS lower bound

Main Question: how large must D be before $\text{val}(\text{SoS}_D) \approx \alpha(G)$?

To prove a lower bound, prove that $\text{val}(\text{SoS}_D) \gg \alpha(G)$ for $D = \text{some } \omega(1)$

□ This is certified by a dual object, which is a degree- D **pseudoexpectation operator** \tilde{E}

Previous techniques for constructing and analyzing \tilde{E} :


- Exact PSD factorization for 3XOR [Grigoriev'99, Schoenebeck'08]
- Symmetric problems [CSS'14, Potechin'18]
- Gram-Schmidt for random, “pairwise uniform” CSPs [KMOW'17]
- Approximate PSD factorization for random, **dense** inputs [BHKMP'16, PR'20, GJJPR'20]
- Lift degree-2 to higher degree [MRX'19, Kunisky'20]

Our contribution: develop lower bound techniques for random, **sparse** inputs

Dense vs Sparse


Dense: uniform $\{-1,+1\}^n$ or $G(n, \frac{1}{2})$ or a collection of Gaussians

- More generally, the input random variables are $O(1)$ -subgaussian



Sparse: $G(n, p)$ or random d -regular graph ($p = d/n$, $d = o(n)$)

- Moments/probability distributions of inputs depend on n
- Our techniques work best for $d \in [\text{polylog}(n), o(n)]$



Independent Set

Given an n -vertex graph G , what is $\alpha(G)$ = largest independent set in G ?

In a random graph $G \sim G(n, d/n)$ with $d = o(n)$, with high probability...

Theorem [Frieze'90]. $\alpha(G) = (2+o_d(1)) n \ln(d) / d$

Theorem [Coja-Oghlan'05]. $\vartheta(G) = \text{val}(\text{SoS}_2) = d^{1/2-o(1)} \alpha(G)$

Theorem [this talk]. Assuming $d \gg \log(n)$, for $D = n^{o(1)}$, $\text{val}(\text{SoS}_D) = d^{1/2-o(1)} \alpha(G)$.

That is, there is a degree- D \tilde{E} satisfying “ G has an ind set of size $d^{1/2-o(1)} \alpha(G)$ ”

Pseudocalibration doesn't work




Boolean Fourier analysis

Recall: Boolean function $f : \{-1, +1\}^n \rightarrow \mathbb{R}$ can be represented in the **Fourier basis**

- This is just its representation as a multilinear polynomial
- There is one monomial/Fourier character χ_S for each subset $S \subseteq [n]$
- Useful when f is a function of a *random* bitstring x , because the basis is orthonormal

In the Independent Set problem, the input is a graph $G \in \{-1, +1\}^{\binom{n}{2}}$

- There is one monomial/Fourier character χ_H for each graph H on $[n]$



Example:

$$\chi_H = G_{1,2} G_{2,3} G_{3,5} G_{4,5}$$

Example:

$$\{1,2,3\} \text{ are a triangle} = (1 - G_{1,2})(1 - G_{1,3})(1 - G_{2,3})/8$$

$$\left(\begin{array}{cccc} \text{graph 1} & - & \text{graph 2} & - & \text{graph 3} \\ \text{graph 4} & + & \text{graph 5} & + & \text{graph 6} \\ \text{graph 7} & + & \text{graph 8} & + & \text{graph 9} \\ \text{graph 10} & - & \text{graph 11} & - & \text{graph 12} \end{array} \right) / 8$$

Pseudocalibration

We want to construct $\tilde{\mathbb{E}}$ that “thinks” there is a large independent set

$\tilde{\mathbb{E}}$ is a *fake* expectation over independent sets, specified by moments up to degree D

$$\tilde{\mathbb{E}}[X^V] \text{ for each } V \subseteq [n], |V| \leq D, \quad \tilde{\mathbb{E}}[X_1] = 0.2, \quad \tilde{\mathbb{E}}[X_7 X_{10}] = 0.75$$

[BHKMP'16] suggests **pseudocalibrating** $\tilde{\mathbb{E}}$ using a **planted distribution** D

Suppose $(G', S) \sim D$ is a distribution on graphs G' with large ind sets

When given the (non-planted) input graph $G \sim G(n, p)$, for $V \subseteq [n]$ define:

$$\tilde{\mathbb{E}}[X^V] = \sum_{H: |V \cup V(H)| \leq \tau} \mathbb{E}_{(G', S) \sim D} [S^V \cdot \chi_H(G')] \cdot \chi_H(G)$$

Pseudocalibration

Natural planted distribution for an Independent Set of size k :

- Sample $G \sim G(n,p)$
- Plant an independent set of expected size k : pick each vertex independently with probability k/n and remove edges between picked vertices

Problem: this distribution is distinguishable from $G(n,p)$ with probability $\Omega(1)$




- Distinguisher: count the number of edges in G
- In all previous successful uses of pseudocalibration, the planted and random distributions are (conjecturally) hard to distinguish by all polynomial time algorithms

Connected Truncation

Solution: truncate away the global tests when defining $\tilde{\mathbf{E}}$


$$\tilde{\mathbf{E}}[X^V] = \sum_{\substack{H: |V \cup V(H)| \leq \tau, \\ H \text{ connected to } V}} \mathbb{E}_{(G', S) \sim D} [S^V \cdot \chi_H(G')] \cdot \chi_H(G)$$

$V = \{1, 2\}$



All that remains is to prove that $\tilde{\mathbf{E}}$ is PSD

Matrix-valued Fourier analysis




Boolean Fourier analysis: matrix edition


We need to analyze matrix-valued functions of the graph G

- Each entry of the matrix is itself a function of G

Def: a **ribbon** is a particular Fourier character in a particular matrix entry. It is specified by a graph H on $[n]$ and two index sets $A, B \subseteq [n]$




$G_{1,2} G_{2,3} G_{3,5} G_{4,5}$



Boolean Fourier analysis: matrix edition

Def: graph matrix M_α sums together all ribbons with a particular shape α



Theorem. Where max is over vertex separators S of the left and right sides of α , whp:

$$\|M_\alpha\| \leq \max_S \tilde{O}(n^{|V(\alpha)|/2 + \text{Isolated}(\alpha)/2} n^{-|V(S)|/2} p^{-|E(S)|/2})$$

Boolean Fourier analysis: matrix edition

The spectral norm of a graph matrix is determined by combinatorial properties of the shape!

Dense setting

Theorem [AMP'16].

$$\|M_\alpha\| \leq \max_S \tilde{O}(n^{|V(\alpha)|/2 + \text{Isolated}(\alpha)/2} n^{-|V(S)|/2})$$

Sparse setting

Theorem [this talk].

$$\|M_\alpha\| \leq \max_S \tilde{O}(n^{|V(\alpha)|/2 + \text{Isolated}(\alpha)/2} n^{-|V(S)|/2} p^{-|E(S)|/2})$$

We use an approximate PSD decomposition [BHKKMP'16].

Vertex factors can be handled in a similar way to [BHKKMP'16].

For edge factors, we give new charging arguments.

Further technical details

- Spectral norms are not concentrated enough due to rare events (presence of small dense subgraphs, e.g. K_8). **Condition** on no such subgraphs.
- Some spectral norms should be controlled using the Frobenius norm
- Matrix for \tilde{E} has a **nullspace**. Factor it out with **missing edge indicators**.
- Use **quasi-missing edge indicators** and missing edge indicators to handle edges incident to $U_\alpha \cap V_\alpha$
- Even with all this work, there is still one log factor in our bound. Removing it would allow us to prove bounds for the regime of constant d



Conclusion

We develop lower bound technology for SoS in the **sparse** setting

- Even though there isn't a planted distribution, we amend pseudocalibration by using a **connected truncation**
- Prove **norm bounds** for sparse graph matrices
- Combinatorial analysis using **conditioning** and **apx PSD factorization**

Open problems:

- Is there a planted distribution which is hard to distinguish from $G(n,p)$?
- Extend Independent Set lower bound to the regime of constant d
- SoS lower bounds for other sparse problems: MaxCut on a random degree- d graph, Densest- k -Subgraph



Thanks for watching!

